

# Mobile Data Management

Dominic Duggan

Based on material by Douglas Terry

1

## **SYSTEM MODELS**

2

## System Models

- Issue: How to make data available to mobile devices
- Approaches:
  - Remote (server-based) data access
  - Device-master replication
  - Peer-to-peer replication
  - Publish-subscribe

3

## Remote Data Access



- Retrieve data from server (always!)
- Connectivity
- Battery life



4


## Device-Master Replication

- Device-side caching
  - Ex: caching in CODA
  - On-demand caching
  - Hoarding/stashing
- Active, user-visible replica
  - Ex: syncing PDA, cell phone calendar
  - Weakened consistency requirements (stale data)
  - Potential for update conflicts



5

## Peer-to-peer replication




- Aka **multi-master replication**
- Any node in overlay network can propagate updates to other nodes
- Advantage: nodes do not need access to server to synchronize data
- Disadvantage:
  - More complex model (no master copy of data)
  - Lack of knowledge about replication topology
  - Decentralized conflict handling

6

# Publish-subscribe


- Publisher emits snippets of data on channels
- Ex: Cellular providers broadcast news items
- Ex: Broadcast disks



7

# Related

- Sensor networks
- Delay-tolerant networking
  - Physical device movement is part of routing
- Infostations
  - Staging area for send/receiving data



8

## Replication Requirements

	Remote Access	Device Master	Peer to Peer	Publish Subscribe
Continuous Connectivity	✓✓			✓
Update Anywhere		✓✓	✓✓	
Consistency		✓✓	✓✓	✓✓
Topology Independence			✓✓	
Conflict Handling		✓✓	✓✓	
Partial Replication		✓	✓	✓✓

## REPLICATED DATA PROTOCOLS

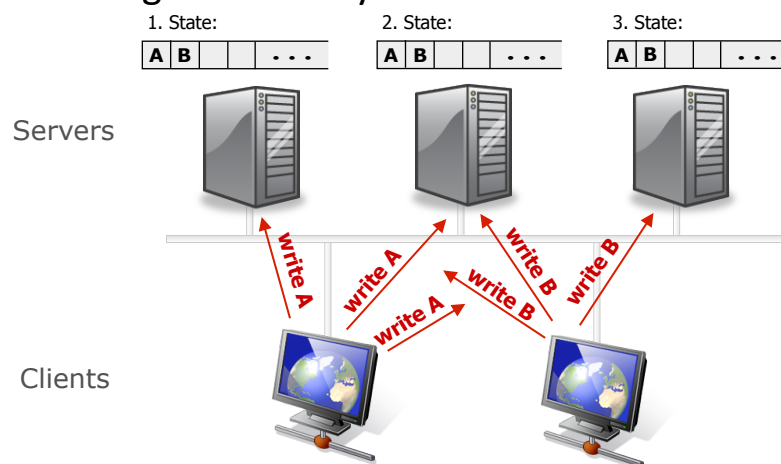
## Questions for Data Replication

- What consistency requirements do we guarantee for replicated data?
- How do we represent the updates?
- How do we send updates?
- How do we order the updates?

11

## Data Consistency


- Strong consistency



12

# Eventual Consistency


- Eventual consistency
  - Two properties:
    - Total propagation
    - Consistent ordering
  - Ordering of updates important



13

## What to do with out-of-order updates?

- Delay delivery
  - Ordered multicast
  - Delay even for local replica!
  - Not practical for disconnected operation
- Tentative update
  - Resolve with other updates later
    - Undo, perform missing updates, redo
  - Information about non-conflicting updates
    - API to specify



14

## Questions for Data Replication

- What consistency requirements do we guarantee for replicated data?
- How do we represent the updates?
- How do we send updates?
- How do we order the updates?

15

## Representing Updates

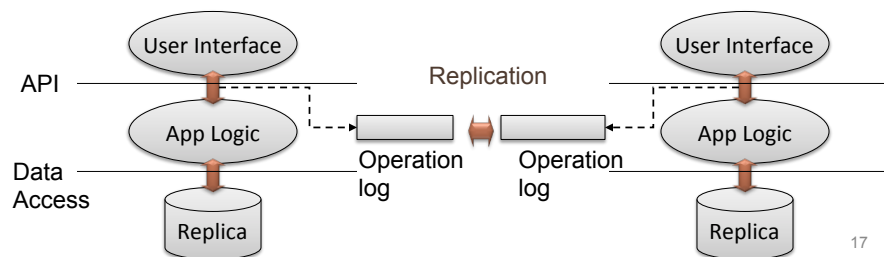
- Operation-sending protocols
  - Ex: File system updates in CODA
- Item-sending protocols
  - Ex: Disk block updates

16



## Recording Updates (1)

- Log-based systems
  - Replicas will converge to consistent state provided:
    - Each replica receives and applies all updates operations
    - Non-commutative operations applied in **same order**
    - Operations have deterministic execution
    - Ex: CODA, Bayou




17

Replica = file system  
Data item = file

## Recording Updates (2)

- State-based systems
  - Attach metadata to **items** (modified bit, update timestamp, version number)
  - Deleted items: create-delete ambiguity
    - “deleted” bit (**tombstone**, **death cert**)
  - Cannot take advantage of operation semantics e.g. to ensure atomicity in update replay




18

## Questions for Data Replication

- What consistency requirements do we guarantee for replicated data?
- How do we represent the updates?
- How do we send updates?
- How do we order the updates?

19


## Gossip



- [t=0] Suppose I know something
- [t=1] I pick you... Now two of us know it.
- [t=2] We each pick ... now 4 know it...
- Information spread: exponential rate.
  - Due to re-infection (gossip to an infected node) spreads as  $1.8^k$  after  $k$  rounds
  - But in  $O(\log(N))$  time,  $N$  nodes are infected

20


## Gossip



21

## Gossip scales very nicely


- Participants' loads independent of size
- Network load linear in system size
- Data spreads in  $\log(\text{system size})$  time



22

# Message Queue Protocols

- Message queue protocols
  - Multicast tree for routing
  - Must notify sender if delivery fails



23

# Modified bit protocol

- Modified bit protocol
  - Simple but only pairwise e.g. hot-syncing PDA




24

## Device-master timestamp protocols

- Device-master timestamp protocols
  - Timestamp or update counter
  - Either record per-client timestamp on server, or clients must record timestamp of last synchronization with master

Replica = file system  
Data item = file



25

## Device-master log-based protocols

- Device-master log-based protocols
  - Update log e.g. CODA, Rover



26

## Full Replica or Data Exchange


- Full replica or data exchange
  - “Flooding” protocol
  - Pairwise between replicas, log or state-based
  - Expensive



27

## Anti-Entropy Protocols

- Anti-Entropy Protocols
  - Peer-to-peer replication
  - Timestamps don't scale with # of peers
    - Each device keeps timestamps of peers for each item
    - Each update communicated by every partner
  - **Meta-data exchange**



28

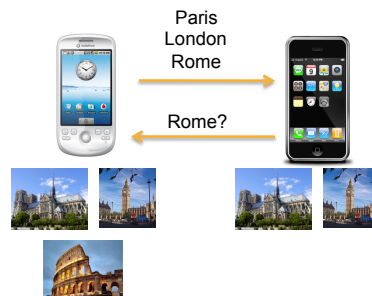
## Anti-Entropy Protocols

- Anti-Entropy Protocols
  - Peer-to-peer replication
  - Timestamps don't scale
  - **Meta-data exchange:**
    - peer shares meta-data




## Anti-Entropy Protocols

- Anti-Entropy Protocols
  - Peer-to-peer replication
  - Timestamps don't scale
  - **Meta-data exchange:**
    - peer shares meta-data
    - peers ask for missing data



## Anti-Entropy Protocols


- Anti-Entropy Protocols
  - Peer-to-peer replication
  - Timestamps don't scale
  - **Meta-data exchange:**
    - peer shares meta-data
    - peers ask for missing data



31

## Anti-Entropy Protocols

- Anti-Entropy Protocols
  - Peer-to-peer replication
  - Timestamps don't scale with # of peers
  - **Meta-data exchange:**
    - peer shares meta-data for replica
    - including version numbers
    - then data exchange
  - **Each update received exactly once**
  - Expensive



32



## Anti-Entropy with checksums

Replica = file system  
Data item = file

- Anti-Entropy with checksums
  - Exchange checksums
  - Then meta-data exchange
  - Then data exchange
- First exchange meta-data for log items
  - Exchange missing log entries before computing checksums
  - **Peel-back check-summing:**
    - Exchange checksums
    - If checksums don't match, exchange latest version of an item
    - Re-compute checksums and iterate

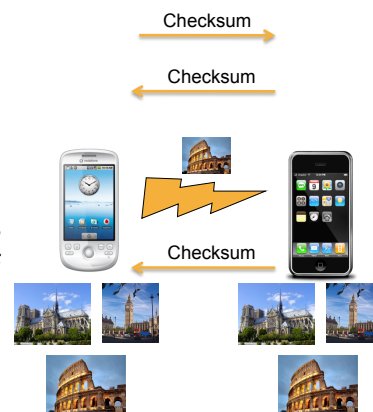


33

## Anti-Entropy with checksums

Replica = file system  
Data item = file

- **Peel-back check-summing:**
  - Exchange checksums
  - If checksums don't match, exchange latest version of an item
  - Re-compute checksums and iterate



34

## Knowledge-driven log-based protocols

- Assume each operation uniquely identified
  - “**knowledge**” of a device: operations it knows about
  - synchronize by sending knowledge
  - store meta-data per replica rather than per data item

Replica = file system  
Data item = file



Source



Target

35

## Knowledge-driven log-based protocols

- **Accept-timestamp** = (replica id, update counter)
  - Uniquely identifies each operation
- **Knowledge vector**: set of accept-timestamps
  - One entry per replica
  - Assume updates received in order they originated



(Joe's reader, 5)



(Joe's laptop, 17)

36

## Knowledge-driven log-based protocols

- Synchronization:

- Target sends its **knowledge vector** to source



- Source sends **logged updates** that are missing



- Updates originating on same device must be kept in order


- Can be resumed if connection is interrupted

37

## Knowledge-driven log-based protocols

- Synchronization:


- Target sends its **knowledge vector** to source



38

## Knowledge-driven log-based protocols

- Synchronization:
  - Source sends **logged updates** that are missing



39

## Knowledge-driven log-based protocols

- Freeing up logs
  - Global protocol
  - Or discard log prefixes, keep omitted vector
    - Vector of knowledge for discarded logs
    - Fail over to metadata-driven protocol if necessary
      - If target knowledge omits updates pruned from source logs

40




## Knowledge-driven **state**-based protocols



- Ex: WinFS
- **Accept-timestamp** = (replica id, update counter)
  - Each data item has a version number
  - AS = (replica id, version) where version is highest version known to have originated at replica
- **Knowledge vector** is set of versions instead of set of operations

41

## Knowledge-driven state-based protocols

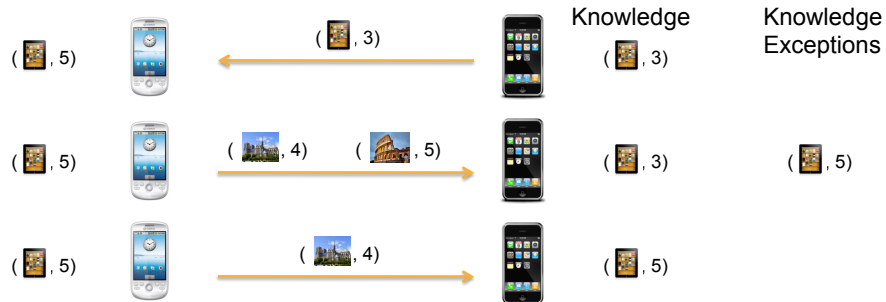
- Knowledge-driven state-based protocols
  - Synchronization:
    - Target sends its **knowledge vector** to source  
 ( , 3)
    - Source sends **versions of items** that are missing on the target  
 ( , 4) ( , 5)
  - Challenge: keeping updates in order
    - No update logs to store updates in order
    - Too expensive to sort at source



42

## Knowledge-driven state-based protocols

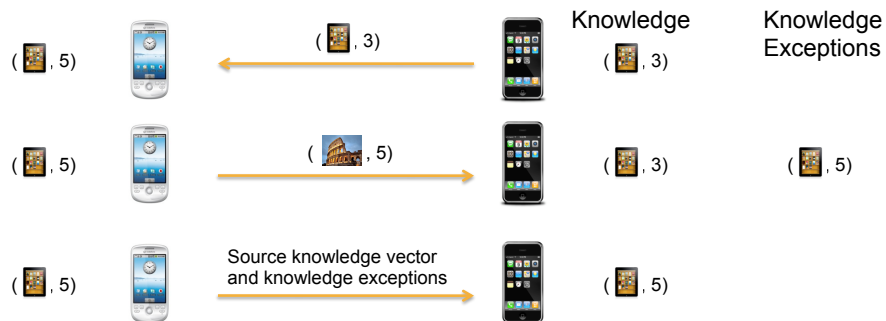
- Knowledge-driven state-based protocols
  - **Knowledge Exception**: for out of order updates



43

## Knowledge-driven state-based protocols

- Knowledge-driven state-based protocols
  - **Learned knowledge**: to remove knowledge exns



44


## Questions for Data Replication

- What consistency requirements do we guarantee for replicated data?
- How do we represent the updates?
- How do we send updates?
- **How do we order the updates?**

45

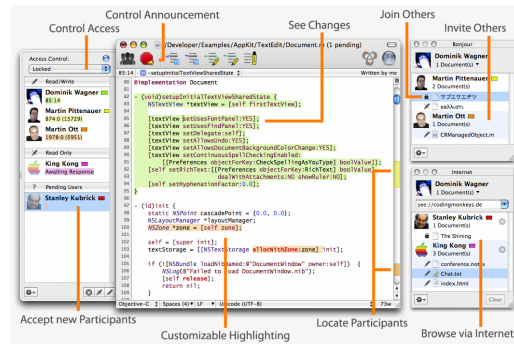
## Ordering Updates (1)

- Ordered delivery
- Sequencers
  - Or Paxos algorithm
- Update timestamps
  - Logical timestamps ensure causal order
- Update counters
  - Concurrent updates
  - Resolve with device id
- Version vectors (like vector clocks)



## Ordering Updates (2)


- Operation transformation
  - Used e.g. for concurrent editing
  - For  $n$  operations, need  $n^2$  transformations



47

## Ordering Updates (2)

- Operation transformation
  - Used e.g. for concurrent editing
  - For  $n$  operations, need  $n^2$  transformations



Windows sucks the

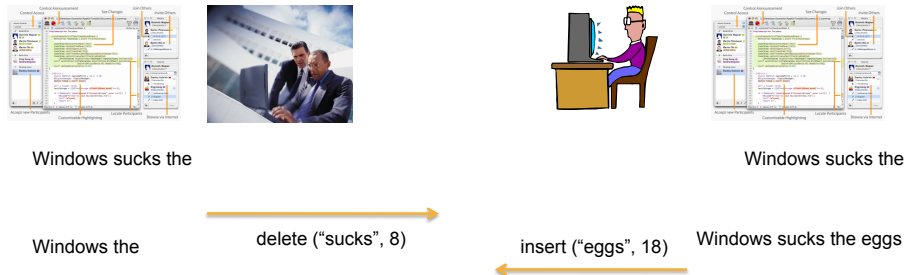


Windows sucks the



## Ordering Updates (2)

- Operation transformation
  - Used e.g. for concurrent editing
  - For  $n$  operations, need  $n^2$  transformations



## Ordering Updates (2)

- Operation transformation
  - Used e.g. for concurrent editing
  - For  $n$  operations, need  $n^2$  transformations



## PARTIAL REPLICATION

51

### Access-based caching

- Access-based caching
  - Problems with callbacks (cf AFS):
    - Callbacks may be missed
    - Server may not know which items are cached
    - Use modified-bit, update timestamp, metadata exchange or knowledge-driven protocols instead
  - Knowledge-driven protocols: devices that are synchronizing are assumed to cache same items
  - Metadata exchange: expensive unless minimize # of items involved
    - sync initiated by caching device



52

## Partial replication

- Policy-based hoarding
  - CODA
- Topic-based channels
  - Disseminate with gossip
  - Peer-to-peer
    - May receive different channels from different sets of peers
    - Knowledge-driven: one knowledge vector per channel

53

## Hierarchical sub-collections


- Hierarchical sub-collections
  - Ex: Mail folders on laptop, inbox only on cell phone
  - Device-master: master only sends updates for sub-collection
  - Peer-to-peer: single knowledge vector will not handle sub-collections



54

## Hierarchical sub-collections

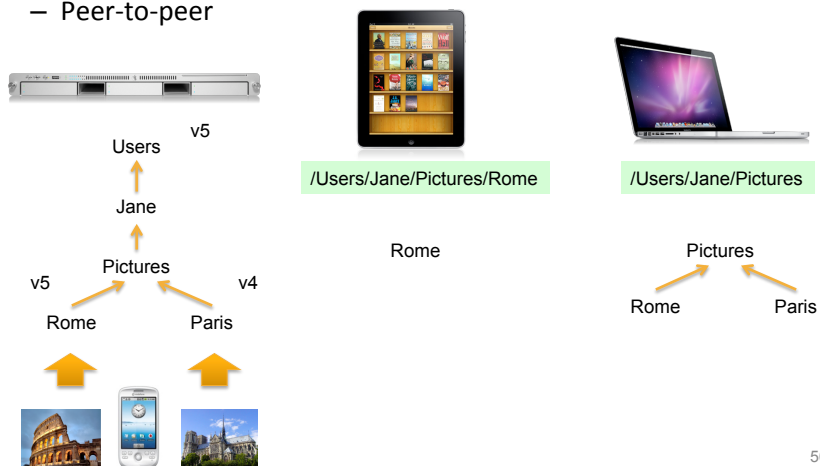
- Hierarchical sub-collections
  - Peer-to-peer



55

## Hierarchical sub-collections

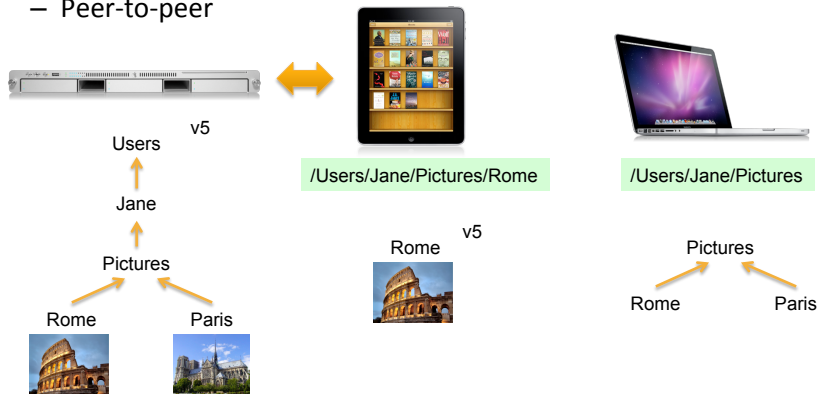
- Hierarchical sub-collections
  - Peer-to-peer



56

## Hierarchical sub-collections

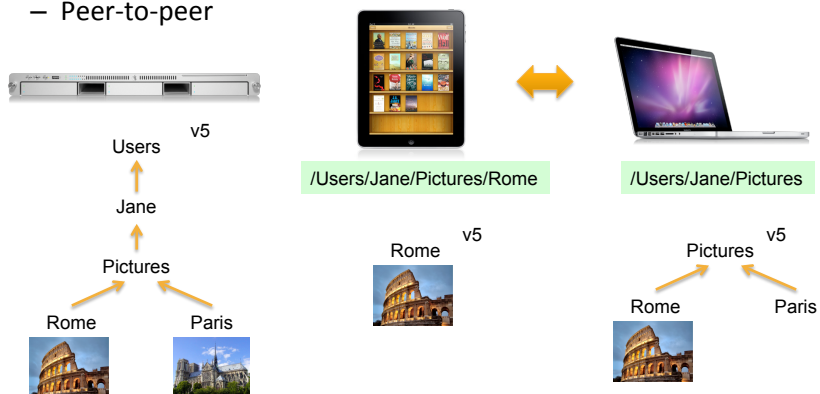
- Hierarchical sub-collections
  - Peer-to-peer



57

## Hierarchical sub-collections

- Hierarchical sub-collections
  - Peer-to-peer



58

## Hierarchical sub-collections

- Hierarchical sub-collections
  - Ex: Mail folders on laptop, inbox only on cell phone
  - Device-master: master only sends updates for sub-collection
  - Peer-to-peer: single knowledge vector will not handle sub-collections
    - Separate knowledge vector for each sub-collection
    - Updates may be received more than once



59

## Content filters

- Content filters
  - Filter query should be on sync partners
  - Device-master state-based: easy
  - Log-based: filter based on type or operation or item it updates
  - **Move-out**: what if cached item is updated and no longer matches filter?
  - Peer-to-peer: topology issue (full-partial-full)
- Metadata exchange supports content filters and move-out
  - open issue for knowledge-driven

60

## Context filters

- Context filter
  - Need access to contextual information
    - Calendar
    - Location information
  - Ex: Cogenia Context Server
  - Move-outs are critical

61

## CONFLICT MANAGEMENT

62


# Conflict Management

- What is a conflict?
  - Write-write; write-read ignored in mobile
  - Single-objects vs multi-object concurrency conflict
  - Transactional conflicts in databases
    - Optimistic concurrency control
    - $\text{ReadSet}(T1) \cap \text{WriteSet}(T2) \neq \{ \}$  and  $\text{WriteSet}(T1) \cap \text{ReadSet}(T2) \neq \{ \}$
  - Operational conflicts
    - Deposits vs withdrawals
  - Semantic (application-specific) conflicts
    - Calendar entries, file names in dir, employee vs manager salary

63

## Conflict Detection (1)

- No conflict detection
  - Grapevine: use timestamps to choose most recent version
    - Danger if a device has a slow clock
- Version histories
  - Digraph: node=version, edge=causal dependency
  - In merged versions, conflict if two versions with no paths between them
  - Expensive



64




## Conflict Detection (2)

- Previous versions
  - Only store the previous version with an item



- Suppose conflict if neither version is predecessor of the other?




## Conflict Detection (2)

- Previous versions
  - Ordered logs will ensure conflicts are found, if each replica sees all updates



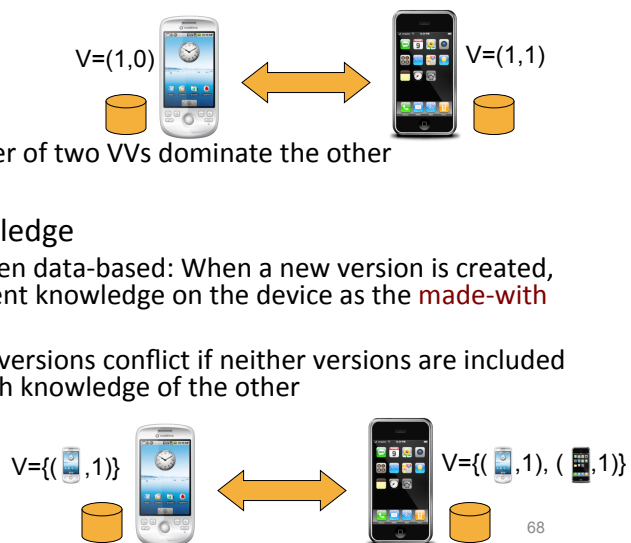
## Conflict Detection (2)

- Previous versions
  - Ordered logs will ensure conflicts are found, if each replica sees all updates



## Conflict Detection (3)

- Version vectors
  - Conflict if neither of two VVs dominate the other
- Made-with Knowledge
  - Knowledge-driven data-based: When a new version is created, record the current knowledge on the device as the **made-with knowledge**
  - Alice and Bob's versions conflict if neither versions are included in the made-with knowledge of the other



## Conflict Detection (4)

- Read-sets
  - Read-sets for optimistic transactions
  - When an item is received at device, does its read-set include an item that has a different version on the device?
- Operation conflict tables
  - Conflict may be based on parameters
  - Requires infinite log to search for conflicting operations
  - Impractical
- Integrity constraints
  - Specify constraints as data invariants
  - Database triggers for violations

69

## Conflict Detection (5)

- Dependency check
  - General scheme for implementing conflict detection
  - For each logged update, store a query and an expected set of results
  - Examples
    - Previous version
    - Read-set
    - Integrity constraints

70

# Conflict Resolution

- How?
  - Manual—conflict log
  - Conflict resolution policy—who
  - Conflict resolvers—dangerous
- Where?
  - **Resolve everywhere**: requires deterministic conflict resolution (Bayou mergeprocs)
  - **Resolve anywhere**: device propagates new updates as a result of automatic resolution
    - Conflict resolution servers
    - Danger of conflict resolution wars
      - Don't use conflict resolver on conflict produced by conflict resolver

71